

Lecture 16:
Multicollinearity
STAT 432, Spring 2021

Multicollinearity

When predictors in a regression model are strongly correlated there can be a number of issues:

- ▶ Large changes in the estimated regression coefficients when a predictor variable is added or dropped.
- ▶ The signs of the coefficients can be the opposite of what we expect.
- ▶ The standard errors are inflated so the t-tests may fail to reveal significant predictors.
- ▶ The predictor variables are not significant when the overall F-test is highly significant.

Multicollinearity

Multicollinearity can be detected in several ways:

- ▶ Examining the relationships between the predictor variables in the scatter plot matrix.
- ▶ Examining the correlation matrix of the predictor variables.
- ▶ Variance inflation factors (VIFs).

Variance Inflation Factor (VIF)

Consider the multiple linear regression model

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p + \epsilon$$

The **variance inflation factor** for predictor x_j is computed as

$$VIF_j = \frac{1}{1 - R_j^2}$$

where R_j^2 is the coefficient of determination obtained from the regression of x_j on all other predictors in the model (i.e., the percent of variation in x_j explained by the other predictors).

As a rough rule, we suspect multicollinearity with predictors for which $VIF_j > 5$, which is equivalent to $R_j^2 > 0.8$.

Variance Inflation Factor (VIF)

For example, suppose that $R_j^2 = 0.99$, then

$$VIF_j = \frac{1}{1 - 0.99} = 100$$

The interpretation is that the standard error of the estimated regression coefficient, $se(\hat{\beta}_j)$, is $\sqrt{100} = 10$ times larger than it would be if there was no collinearity ($R_j^2 = 0$).

Example¹


- ▶ Car drivers like to adjust the seat position for their own comfort. Car designers would find it helpful to know where different drivers will position the seat depending on their size and age.
- ▶ Researchers at the HuMoSim laboratory at the University of Michigan collected data on 38 drivers.
- ▶ The response variable is `hipcenter`, the horizontal distance of the midpoint of the hips from a fixed location in the car in mm.
- ▶ The predictors of interest are age, weight, height with and without shoes, seated height, arm length, thigh length, and lower leg length.

¹Example from Julian Faraway, *Linear Models in R*, 1st edition, Ch. 5, pp. 83–87

```
> library(faraway)
```

```
> head(seatpos)
```

	Age	Weight	HtShoes	Ht	Seated	Arm	Thigh	Leg	hipcenter
1	46	180	187.2	184.9	95.2	36.1	45.3	41.3	-206.300
2	31	175	167.5	165.5	83.8	32.9	36.5	35.9	-178.210
3	23	100	153.6	152.2	82.9	26.0	36.6	31.0	-71.673
4	19	185	190.3	187.4	97.3	37.4	44.1	41.0	-257.720
5	23	159	178.0	174.1	93.9	29.5	40.1	36.9	-173.230
6	47	170	178.7	177.0	92.4	36.0	43.2	37.4	-185.150



There are several strong correlations between the predictors.

```
> round(cor(seatpos[, -9]), 2)
```

	Age	Weight	HtShoes	Ht	Seated	Arm	Thigh	Leg
Age	1.00	0.08	-0.08	-0.09	-0.17	0.36	0.09	-0.04
Weight	0.08	1.00	0.83	0.83	0.78	0.70	0.57	0.78
HtShoes	-0.08	0.83	1.00	1.00	0.93	0.75	0.72	0.91
Ht	-0.09	0.83	1.00	1.00	0.93	0.75	0.73	0.91
Seated	-0.17	0.78	0.93	0.93	1.00	0.63	0.61	0.81
Arm	0.36	0.70	0.75	0.75	0.63	1.00	0.67	0.75
Thigh	0.09	0.57	0.72	0.73	0.61	0.67	1.00	0.65
Leg	-0.04	0.78	0.91	0.91	0.81	0.75	0.65	1.00

The model shows signs of multicollinearity. The overall F statistic is large and $R^2 = 0.6$, but none of the individual predictors are significant.

```
> lm1 <- lm(hipcenter ~ ., data=seatpos)
```

```
> summary(lm1)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	436.43213	166.57162	2.620	0.0138 *
Age	0.77572	0.57033	1.360	0.1843
Weight	0.02631	0.33097	0.080	0.9372
HtShoes	-2.69241	9.75304	-0.276	0.7845
Ht	0.60134	10.12987	0.059	0.9531
Seated	0.53375	3.76189	0.142	0.8882
Arm	-1.32807	3.90020	-0.341	0.7359
Thigh	-1.14312	2.66002	-0.430	0.6706
Leg	-6.43905	4.71386	-1.366	0.1824

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 37.72 on 29 degrees of freedom

Multiple R-squared: 0.6866, Adjusted R-squared: 0.6001

F-statistic: 7.94 on 8 and 29 DF, p-value: 1.306e-05

Several of the variance inflation factors are large and exceed the 5 cut-off.

```
> library(faraway) # to use vif() function
```

```
> round(vif(lm1), 2)
```

Age	Weight	HtShoes	Ht	Seated	Arm	Thigh	Leg
2.00	3.65	307.43	333.14	8.95	4.50	2.76	6.69

For HtShoes the interpretation is that the standard error for this predictor is $\sqrt{307.4} = 17.5$ times larger than it would be without collinearity.

We can also compute the VIFs manually.

```
# create data frame only containing predictors
```

```
> x <- seatpos[, -9]
```

```
> summary(lm(Ht ~., data=x))$r.squared
```

```
[1] 0.9969982
```

```
> 1/(1 - 0.9969982)
```

```
[1] 333.1335
```

Your Turn

Manually compute the VIF for the predictor Seated, which is seated height in cm.

- ▶ Many of the variables in the full model are redundant, and do the same job at predicting the response.
- ▶ For example, the following predictors all measure the length of the driver in some way: `HtShoes`, height in shoes; `Ht`, height bare foot; `Seated`, seated height; `Arm`, arm length; `Thigh`, thigh length; and `Leg`, leg length.
- ▶ Instead of using all of these length predictors, we can just select one to include in the model, and drop the others.
- ▶ However, because of collinearity, we should not conclude that the variables we drop have nothing to do with the response.

Consider the correlation matrix with just the length variables. All of these predictor variables are strongly correlated with each other. We pick Ht since it is the simplest measure, and more strongly correlated with the response than the other predictors.

```
> round(cor(seatpos[, 3:9]), 2)
```

	HtShoes	Ht	Seated	Arm	Thigh	Leg	hipcenter
HtShoes	1.00	1.00	0.93	0.75	0.72	0.91	-0.80
Ht	1.00	1.00	0.93	0.75	0.73	0.91	-0.80
Seated	0.93	0.93	1.00	0.63	0.61	0.81	-0.73
Arm	0.75	0.75	0.63	1.00	0.67	0.75	-0.59
Thigh	0.72	0.73	0.61	0.67	1.00	0.65	-0.59
Leg	0.91	0.91	0.81	0.75	0.65	1.00	-0.79
hipcenter	-0.80	-0.80	-0.73	-0.59	-0.59	-0.79	1.00

Removing some correlated predictors fixes many of the issues caused by multicollinearity. The predictor Ht is now highly significant in the model. Further simplification is clearly possible.

```
> lm2 <- lm(hipcenter ~ Age + Weight + Ht, data=seatpos)
> summary(lm2)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	528.297729	135.312947	3.904	0.000426	***
Age	0.519504	0.408039	1.273	0.211593	
Weight	0.004271	0.311720	0.014	0.989149	
Ht	-4.211905	0.999056	-4.216	0.000174	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 36.49 on 34 degrees of freedom

Multiple R-squared: 0.6562, Adjusted R-squared: 0.6258

F-statistic: 21.63 on 3 and 34 DF, p-value: 5.125e-08

```
> vif(lm2)
```

	Age	Weight	Ht
	1.093018	3.457681	3.463303

The R^2 of the reduced model with just 3 predictors (Age, Weight, and Ht) is close to the R^2 of the full model with all the strongly correlated predictors. In fact, the adjusted R^2 for the reduced model is slightly higher than the full model.

```
> summary(lm1)$r.squared
```

```
[1] 0.6865535
```

```
> summary(lm2)$r.squared
```

```
[1] 0.6561654
```

```
> summary(lm1)$adj.r.squared
```

```
[1] 0.6000855
```

```
> summary(lm2)$adj.r.squared
```

```
[1] 0.6258271
```

Concluding Remarks

Some ways to deal with multicollinearity:

- ▶ If several predictors are strongly correlated with each other, pick one predictor out of the bunch to use in the reduced model. The R^2 should not change much after removing some correlated predictors.
- ▶ You can also combine predictors. For instance, by taking the sum or average of two correlated predictors.
- ▶ Automated stepwise variable selection techniques can also be used to find an optimal subset set of predictors, and thereby reduce collinearity problems.